Wstęp do Uczenia Maszynowego 2020: projekt I, kamień milowy III - Regresja Logistyczna

Marcin Łukaszyk

April 18, 2020

Wczytanie Danych

```
Reg_train <- read.csv("Reg_train.csv")
Reg_test <- read.csv("Reg_test.csv")
Reg_train <- select(Reg_train,-X)
Reg_test <- select(Reg_test,-X)

knitr::kable(sample_n(Reg_test, 5))
```

duration	credit_amount	installment_rate	present_residence	age	existing_credits	dependents	has_telephone
28	7824	3	4	40	2	2	1
24	1597	4	4	54	2	2	0
4	1455	2	1	42	3	2	0
48	4844	3	2	33	1	1	1
24	4057	3	3	43	1	1	1

```
knitr::kable(sample_n(Reg_train, 5))
```

duration	credit_amount	installment_rate	present_residence	age	existing_credits	dependents	has_telephone
24	3512	2	3	38	2	1	1
18	1984	4	4	47	2	1	0
9	918	4	1	30	1	1	0
6	1198	4	4	35	1	1	0
12	1858	4	1	22	1	1	0

WSzystko jest wczytane poprawnie.

Stowrzenie modelu

Naszym modelem będzię model z pakietu scidb. Nie musismy go tworzyć, od razu można "fit'ować" dane do modelu

```r
glm.fit <- glm(is_good_customer_type ~ duration + age + existing_credits + dependents + has_telephone +
               ,data = Reg_train,family = binomial)
```

Parametry modelu

Poniżej znajduje się podsumowanie naszego modelu. Pokazane są wszystkie jego parametry oraz znaczenie w działaniu naszego modelu.

```r
summary(glm.fit)
```

```
##
## Call:
## glm(formula = is_good_customer_type ~ duration + age + existing_credits +
##     dependents + has_telephone + is_foreign_worker + has_problems_credit_history +
##     purpose_domestic + purpose_retraining + purpose_radio_television +
##     purpose_new_car + purpose_used_car + purpose_business + purpose_repairs +
##     purpose_education + purpose_furniture_equipment + other_debtors_guarantor +
##     other_debtors_co_applicant + other_installment_plans_bank +
##     other_installment_plans_stores + housing_rent + housing_own +
##     job_skilled_employee + job_unskilled_resident + job_highly_qualified_employee +
##     savings + present_employment + property + checking_account_status +
##     is_woman + is_single, family = binomial, data = Reg_train)
##
## Deviance Residuals:
##     Min       1Q   Median       3Q      Max
## -2.4722  -0.8899   0.4654   0.7781   2.3977
##
## Coefficients:
##                                 Estimate Std. Error z value Pr(>|z|)
## (Intercept)                     3.384690   1.241297   2.727  0.00640 **
## duration                       -0.042163   0.007819  -5.392 6.97e-08 ***
## age                             0.012578   0.009399   1.338  0.18084
## existing_credits               -0.394848   0.188578  -2.094  0.03628 *
## dependents                     -0.486423   0.265525  -1.832  0.06696 .
## has_telephone                   0.456194   0.207464   2.199  0.02788 *
## is_foreign_worker              -1.482214   0.768042  -1.930  0.05362 .
## has_problems_credit_history     0.782436   0.226453   3.455  0.00055 ***
## purpose_domestic                0.302664   0.334697   0.904  0.36584
## purpose_retraining             -0.772425   0.448215  -1.723  0.08483 .
## purpose_radio_television       -0.042468   0.352021  -0.121  0.90398
## purpose_new_car                -0.671123   0.335964  -1.998  0.04576 *
## purpose_used_car                1.109216   0.451582   2.456  0.01404 *
## purpose_business                1.009181   1.142224   0.884  0.37695
## purpose_repairs                -0.334217   0.843957  -0.396  0.69210
## purpose_education              -0.543683   0.601829  -0.903  0.36632
## purpose_furniture_equipment     0.439982   0.828232   0.531  0.59526
## other_debtors_guarantor         0.634619   0.451470   1.406  0.15982
## other_debtors_co_applicant     -0.811476   0.444269  -1.827  0.06777 .
## other_installment_plans_bank   -0.367892   0.248429  -1.481  0.13864
## other_installment_plans_stores -0.702364   0.387586  -1.812  0.06996 .
## housing_rent                   -0.143903   0.389959  -0.369  0.71211
## housing_own                     0.270075   0.336478   0.803  0.42218
```

```
## job_skilled_employee           -0.647848   0.638322  -1.015  0.31014
## job_unskilled_resident         -0.674382   0.655113  -1.029  0.30329
## job_highly_qualified_employee  -1.184181   0.672561  -1.761  0.07829 .
## savings                         0.167907   0.098253   1.709  0.08746 .
## present_employment              0.084757   0.040405   2.098  0.03593 *
## property                        0.174283   0.105706   1.649  0.09920 .
## checking_account_status        -0.429361   0.093033  -4.615 3.93e-06 ***
## is_woman                        0.217902   0.276982   0.787  0.43146
## is_single                       0.802860   0.276274   2.906  0.00366 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##     Null deviance: 975.68  on 799  degrees of freedom
## Residual deviance: 795.81  on 768  degrees of freedom
## AIC: 859.81
##
## Number of Fisher Scoring iterations: 5
```

Test modelu

Podstawowym parametrem jest stosunek poprawnych odpowiedzi. Tzn jest to prosta średnia z 1 jeśli odpowiedź jest dobra i 0 w przeciwnym przypadku.

```
glm.probs <- ifelse(predict(glm.fit,newdata = Reg_test,type = "response") > 0.5,1,0)
mean(glm.probs == select(Reg_test,is_good_customer_type))
```

```
## [1] 0.7
```

Jak widzimy model średnio dobrze przewiduje 70 % odpowiedzi.

Dokładne zmierzenie jakości modelu

Funkcje pomocnicze które określą nam jakość modelu

```
confusion_matrix_values <- function(confusion_matrix){
  TP <- confusion_matrix[2,2]
  TN <- confusion_matrix[1,1]
  FP <- confusion_matrix[1,2]
  FN <- confusion_matrix[2,1]
  return (c(TP, TN, FP, FN))
}

accuracy <- function(confusion_matrix){
  conf_matrix <- confusion_matrix_values(confusion_matrix)
  return((conf_matrix[1] + conf_matrix[2]) / (conf_matrix[1] + conf_matrix[2] + conf_matrix[3] + conf_ma
}

precision <- function(confusion_matrix){
  conf_matrix <- confusion_matrix_values(confusion_matrix)
```

```r
  return(conf_matrix[1]/ (conf_matrix[1] + conf_matrix[3]))
}

recall <- function(confusion_matrix){
  conf_matrix <- confusion_matrix_values(confusion_matrix)
  return(conf_matrix[1] / (conf_matrix[1] + conf_matrix[4]))
}

f1 <- function(confusion_matrix){
  conf_matrix <- confusion_matrix_values(confusion_matrix)
  rec <- recall(confusion_matrix)
  prec <- precision(confusion_matrix)
  return(2 * (rec * prec) / (rec + prec))
}
```

```r
confusion_matrix_primitive <- table(
  Truth = select(Reg_test,is_good_customer_type)[,1],
  Prediction = glm.probs
  )
knitr::kable(confusion_matrix_primitive)
```

	0	1
0	16	45
1	15	124

```r
accuracy_primitive <- accuracy(confusion_matrix_primitive)
precision_primitive <- precision(confusion_matrix_primitive)
recall_primitive <- recall(confusion_matrix_primitive)
f1_primitive <- f1(confusion_matrix_primitive)

classification_report_primitive <- data.frame(accuracy_primitive, precision_primitive,
                                              recall_primitive, f1_primitive)
colnames(classification_report_primitive) <- c("accuracy", "precision",
                                               "recall", "f1")
knitr::kable(classification_report_primitive)
```

accuracy	precision	recall	f1
0.7	0.7337278	0.8920863	0.8051948

Wyrzucenie Mało Znaczących Zmiennych

Do zwiększenia dokladnośći modelu sprobujemy usunąć ze zmiennych te ktore, według funckji summary(), najmniej wplywaja na nasz model.

```r
glm.fit <- glm(is_good_customer_type ~age + dependents + is_foreign_worker + purpose_domestic + purpose_
               ,data = Reg_train,family = binomial)
glm.probs <- ifelse(predict(glm.fit,newdata = Reg_test,type = "response") > 0.5,1,0)
```

```r
confusion_matrix_primitive <- table(
  Truth = select(Reg_test,is_good_customer_type)[,1],
  Prediction = glm.probs
  )
knitr::kable(confusion_matrix_primitive)
```

	0	1
0	9	52
1	15	124

```r
accuracy_primitive <- accuracy(confusion_matrix_primitive)
precision_primitive <- precision(confusion_matrix_primitive)
recall_primitive <- recall(confusion_matrix_primitive)
f1_primitive <- f1(confusion_matrix_primitive)

classification_report_primitive <- data.frame(accuracy_primitive, precision_primitive,
                                              recall_primitive, f1_primitive)
colnames(classification_report_primitive) <- c("accuracy", "precision",
                                               "recall", "f1")
knitr::kable(classification_report_primitive)
```

accuracy	precision	recall	f1
0.665	0.7045455	0.8920863	0.7873016

Jak widać nie uzyskujemy lepszych rezultatów, a nasze wyniki są nawet lekko gorsze. Spróbujmy usunąć jeszcze kilka najmniej istotnych parametrów, na podstaiwe wskazań funckcji summary()

```r
summary(glm.fit)
```

```
##
## Call:
## glm(formula = is_good_customer_type ~ age + dependents + is_foreign_worker +
##     purpose_domestic + purpose_retraining + purpose_radio_television +
##     purpose_business + purpose_repairs + purpose_education +
##     purpose_furniture_equipment + other_debtors_guarantor + other_debtors_co_applicant +
##     other_installment_plans_bank + other_installment_plans_stores +
##     housing_rent + housing_own + job_skilled_employee + job_unskilled_resident +
##     job_highly_qualified_employee + savings + present_employment +
##     property + checking_account_status + is_woman + is_single,
##     family = binomial, data = Reg_train)
##
## Deviance Residuals:
##     Min       1Q   Median       3Q      Max
## -2.3490  -1.1194   0.5603   0.8502   1.8128
##
## Coefficients:
##                                Estimate Std. Error z value Pr(>|z|)
## (Intercept)                    1.868579   1.122322   1.665   0.0959 .
```

```
## age                               0.020000   0.008831   2.265   0.0235 *
## dependents                       -0.451985   0.251674  -1.796   0.0725 .
## is_foreign_worker                -1.635407   0.765932  -2.135   0.0327 *
## purpose_domestic                  0.468442   0.218721   2.142   0.0322 *
## purpose_retraining               -0.370521   0.351804  -1.053   0.2922
## purpose_radio_television          0.211954   0.237065   0.894   0.3713
## purpose_business                  1.253931   1.090908   1.149   0.2504
## purpose_repairs                   0.150804   0.763680   0.197   0.8435
## purpose_education                -0.448609   0.530250  -0.846   0.3975
## purpose_furniture_equipment       0.522847   0.730910   0.715   0.4744
## other_debtors_guarantor           0.404226   0.432311   0.935   0.3498
## other_debtors_co_applicant       -0.910258   0.425450  -2.140   0.0324 *
## other_installment_plans_bank     -0.431306   0.235634  -1.830   0.0672 .
## other_installment_plans_stores -0.740091   0.379054  -1.952   0.0509 .
## housing_rent                      0.021814   0.363947   0.060   0.9522
## housing_own                       0.342489   0.310810   1.102   0.2705
## job_skilled_employee             -0.537390   0.604666  -0.889   0.3741
## job_unskilled_resident           -0.600522   0.622727  -0.964   0.3349
## job_highly_qualified_employee   -0.742295   0.625014  -1.188   0.2350
## savings                           0.177937   0.092682   1.920   0.0549 .
## present_employment                0.085908   0.038128   2.253   0.0242 *
## property                          0.233094   0.098987   2.355   0.0185 *
## checking_account_status          -0.468764   0.087320  -5.368 7.95e-08 ***
## is_woman                          0.106457   0.265636   0.401   0.6886
## is_single                         0.660575   0.264440   2.498   0.0125 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##     Null deviance: 975.68  on 799  degrees of freedom
## Residual deviance: 862.00  on 774  degrees of freedom
## AIC: 914
##
## Number of Fisher Scoring iterations: 5
```

Te zmienne to:

- age
- is_foreign_worker
- present_employment
- property
- checking_account_status
- is_single

```r
glm.fit <- glm(is_good_customer_type ~ dependents + purpose_domestic + purpose_retraining + purpose_radi
               ,data = Reg_train,family = binomial)
glm.probs <- ifelse(predict(glm.fit,newdata = Reg_test,type = "response") > 0.5,1,0)
```

```
confusion_matrix_primitive <- table(
  Truth = select(Reg_test,is_good_customer_type)[,1],
  Prediction = glm.probs
  )
knitr::kable(confusion_matrix_primitive)
```

	0	1
0	4	57
1	5	134

```
accuracy_primitive <- accuracy(confusion_matrix_primitive)
precision_primitive <- precision(confusion_matrix_primitive)
recall_primitive <- recall(confusion_matrix_primitive)
f1_primitive <- f1(confusion_matrix_primitive)

classification_report_primitive <- data.frame(accuracy_primitive, precision_primitive,
                                              recall_primitive, f1_primitive)
colnames(classification_report_primitive) <- c("accuracy", "precision",
                                               "recall", "f1")
knitr::kable(classification_report_primitive)
```

accuracy	precision	recall	f1
0.69	0.7015707	0.9640288	0.8121212

Jak widać otrzymujemy lepsze wyniki niż poprzednio. Ostatnią metodą niech będzie stworzenie modelu z zmiennych mających największe znaczenie w naszym pierwszym modelu. Pięć zmiennych z największym znaczeniem to :

- purpose_radio_television : **0.90**
- housing_rent : **0.71**
- purpose_repairs : **0.69**
- purpose_furniture_equipment : **0.59**
- housing_own : **0.42**

Stwórzmy teraz model na podstawie

```
glm.fit <- glm(is_good_customer_type ~ purpose_radio_television + purpose_repairs + purpose_furniture_e
               ,data = Reg_train,family = binomial)
glm.probs <- ifelse(predict(glm.fit,newdata = Reg_test,type = "response") > 0.5,1,0)
```

```
confusion_matrix_primitive <- table(
  Truth = select(Reg_test,is_good_customer_type)[,1],
  Prediction = glm.probs
  )
knitr::kable(confusion_matrix_primitive)
```

	1
0	61
1	139

Co ciekawe ten model nie przewidział żadngo 0 czyli złego klienta. Ten model uzyskuje celność **0.65** co jest podobnym wynikiem do reszty. Jednak z powodu nie przewidzenia złych klientów nie można obliczyć reszty statystyk.

Podumowanie

Model uzyskuje podobne parapetry dla różnych zmiennych. Najgorzej wypadł model bez pięciu najmienj istotynych zmiennych. Może to być jednak spowodowane ilosćią danych jak i ich arbitralnym podziałem na zbiór testowy i treninigowy.